

# Computing with ICS-ACI

Ben Seiyon Lee

Pennsylvania State University

Department of Statistics

Graduate Workshop

April 15, 2019

# Outline

- ① High Performance Computing
- ② Accessing ACI + Details
- ③ Transferring Files to/from Local Machines
- ④ Submitting Jobs on ACI
- ⑤ Monitoring Jobs
- ⑥ Parallelization
- ⑦ Exercises

# ACI-ICS High Performance Computing Environment

The Institute for CyberScience Advanced CyberInfrastructure (ICS-ACI) is Penn States's high-performance computing (HPC) infrastructure.

## High Performance Computing

- Provides researchers with the hardware, software, and technical expertise needed to solve problems.
- 1000+ servers with 23,000+ cores and 6 Petabytes of storage
- Appropriate for computationally burdensome tasks (parallelization, storage, memory, long run times)
- Access provided through allocations (department or research lab) or the Open queue

"Getting Started with ACI" is a fantastic resource!

# Accessing ACI-B for Batch Jobs

## Sign up for an account:

- [ICS-ACI Account Sign-up](#)
- [2-Factor Authentication](#)

## Mac

- Open Terminal
- ssh into ACI: `ssh <username>@aci-b.aci.ics.psu.edu`
- Complete 2 Factor Authentication

## Windows

- Open Putty
- Enter `aci-b.aci.ics.psu.edu` in the Host Name field
- Select SSH then X11 and Enable X11 forwarding
- Select Connection then Data and enter your username in the Auto-login username field

# How to run R scripts on ACI

For code development and testing (simple tasks), use the interactive shell ([rstudio.aci.ics.psu.edu](https://rstudio.aci.ics.psu.edu)).

\*I prefer testing on my laptop and syncing files via Github.

For computationally demanding tasks, send batch jobs to ACI-B.

## ACI-B Server Types:

- **Basic:** 2.2 GHz Intel Xeon Processor, 24 CPU/node, 128 GB RAM
- **Standard:** 2.8 GHz Intel Xeon Processor, 20 CPU/node, 256 GB RAM
- **High:** 2.2 GHz Intel Xeon Processor, 40 CPU/node, 1 TB RAM
- **GPU:** 2.5 GHz Intel Xeon Processor, 2 Nvidia Tesla K80 computing modules/server, 24 CPU/server

## 2 Options: SCP client or the command line

### SCP Client

- Cyberduck (Mac) or WinSCP (Windows) or FileZilla (All OS)
- **Host name: datamgr.aci.ics.psu.edu**  
Dedicated hostname for file transfers.
- File Protocol: SFTP
- PSU username and password
- Port 22
- Go through 2 Factor Authentication
- Easy-to-use interface similar to Windows Explorer or Finder.

# Transferring Files - Command Line

## Command line

- Helpful for transferring many files
- Use the dedicated data manager hostname **datamgr.aci.ics.psu.edu**
- Other option: git for small files and the command line for large files

## Examples:

- **Copy a file from local machine to ACI-B:**

```
scp <filename> <username>@datamgr.aci.ics.psu.edu:<ACI  
Directory>
```

```
scp test.txt skl5261@datamgr.aci.ics.psu.edu:~/work/Workshop/
```

- **Copy a file from ACI-B to local machine**

```
scp <username>@datamgr.aci.ics.psu.edu:<filename> <Local  
Directory>
```

```
scp skl5261@datamgr.aci.ics.psu.edu:~/work/Workshop/test.txt  
~/home/Workshop/
```

# Sending Jobs via PBS (Portable Batch System)

## Batch jobs

- Sends jobs to batch scheduler using a **.PBS** file
- Scheduler allocates memory, nodes, and processors as necessary

## Example PBS File:

```
#!/bin/bash
#PBS -A open Allocation
#PBS -l nodes=1:ppn=3 Nodes and processors per node
#PBS -l walltime=72:00:00 Requested wall time
#PBS -l pmem=5GB Memory per processor
#PBS -N ProjName Job Name
#PBS -j oe Prints log file for completion and errors
#PBS -m abe Sends email when job aborts, begins, and ends
#PBS -M psuid@psu.edu Email Address
Rscript test.R Runs the R file
```



# Running and Monitoring Jobs

**Use the command line to submit, monitor, or cancel your jobs.**

## **Submit a Job**

- `qsub <PBSfile>`
- Details should be provided in the PBS file

## **Monitor a Job**

- Show queued or running jobs: `qstat`
- Jobs by userid: `qstat -u <userid>`
- Jobs by jobid: `qstat -i <jobid>`
- Show all jobs: `showq`

## **Cancel a Job**

- `qdel jobid`

## **Conditional Execution**

- Run a job after another job has finished  
`qsub -W depend=afterok:<jobid> <second PBS file>`

## Types of Clusters:

- PSOCK: Small-scale computing environments (1 node/ $<20$  cores)
- MPI: Larger-scale computing environments (2+ nodes/20+ cores)

## USE MPI FOR JOBS REQUIRING 20+ CORES!

- Open allocation: Limited to 100 processors (5 nodes) and 48 hours of walltime.

## R packages for parallelization:

- Snowfall
- Rmpi
- doParallel

# Try it out!

## Instructions:

1. Download and unzip the folder **workshop**
2. Login into ACI
3. Copy the “Workshop” folder to your ACI work directory (~ /work/.)
4. Follow instructions in the Readme file

## Exercises:

- ① **Simple Example run on command line**
- ② **Simple Script using PBS Scheduler + 1 Jobs**
- ③ **Parallel Script using PBS Scheduler**
- ④ **Parallel Script using PBS Scheduler + MPI**
- ⑤ **Simple Script using PBS Scheduler + 10 Jobs**

## Exercises:

- ❶ **Simple Example run on command line**  
~100 seconds (1.7 minutes)
- ❷ **Simple Script using PBS Scheduler + 1 Jobs**  
– This takes ~2 minutes
- ❸ **Parallel Script using PBS Scheduler**  
– This takes ~15 seconds. Major speedup!
- ❹ **Parallel Script using PBS Scheduler + MPI**  
– Limited to 100 processors (5 nodes) on the Open queue.  
– Requires installing Rmpi which can be tricky. Installation script included in exercise files.  
– This takes ~49 seconds. Longer times due to communication costs.
- ❺ **Simple Script using PBS Scheduler + 10 Jobs**  
– This takes ~15 seconds per run. But jobs start at different times.

# The End

# Unix Commands

- Change directories: `cd`
  - Home Directory: `cd`
  - Here: `cd .`
  - Up one directory: `cd ..`
  - All files in the directory: `ls *`
  - Wildcards: `Test* . *.png`
  - Send output to another command (piping): `|`
  - Write command output to a file: `ls > log.txt`

- Create Directory: `mkdir`

```
cd ~/work  
mkdir Workshop
```

- Remove Directory: `rm -r <directory>`

```
rmdir WorkshopB  
ls  
mkdir WorkshopB
```

# Unix Commands

- Move Files: mv

```
mv file1.txt ./WorkshopB/
```

```
mv ../WorkshopB/file1.txt ./WorkshopB/file2.txt
```

- Copy Files: cp

```
cp ../WorkshopB/file1.txt ../WorkshopB/file2.txt
```

- Remove Files: rm

```
rm file1.txt
```

```
rm -r WorkshopB
```

- Access Manual for commands: man

```
man rm
```

```
q
```

- List files: ls

```
ls
```

```
ls ~/work/Workshop
```

# Unix Commands

- Print the current directory:

```
pwd
```

- Past commands:

```
history
```

- Manage permissions for a file:

```
chmod u=rwx,g=rwx,o=rwx file1.txt  
chmod 777 file1.txt
```

The digits represent the permissions for the user, group, and others, in that order. Each digit is a combination of the numbers 4, 2, 1, and 0: 4 stands for "read", 2 stands for "write", 1 stands for "execute", and 0 stands for "no permission." 7 is the combination of permissions 4+2+1 (read, write, and execute), 5 is 4+0+1 (read, no write, and execute), and 4 is 4+0+0 (read, no write, and no execute).